

WELCOM TO OUR HEART DISEASE PREDICTION PROJECT

Milestone 1: Data Collection, Exploration, and Preprocessing

1.Data collection:

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer

# Set visualization settings
sns.set_style('whitegrid') # Set seaborn style
plt.rcParams['figure.figsize'] = (10, 6) # Set default figure size

# 1. Load the dataset
df = pd.read_csv('heartT.csv')
df.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
slope \										
0	52	1	0	125	212	0	1	168	0	1.0
2										
1	53	1	0	140	203	1	0	155	1	3.1
0										
2	70	1	0	145	174	0	1	125	1	2.6
0										
3	61	1	0	148	203	0	1	161	0	0.0
2										
4	62	0	0	138	294	1	1	106	0	1.9
1										

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

2.Data Exploration

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1025 entries, 0 to 1024
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	age	1025 non-null	int64
1	sex	1025 non-null	int64
2	cp	1025 non-null	int64
3	trestbps	1025 non-null	int64
4	chol	1025 non-null	int64
5	fbs	1025 non-null	int64
6	restecg	1025 non-null	int64
7	thalach	1025 non-null	int64
8	exang	1025 non-null	int64
9	oldpeak	1025 non-null	float64
10	slope	1025 non-null	int64
11	ca	1025 non-null	int64
12	thal	1025 non-null	int64
13	target	1025 non-null	int64

```
dtypes: float64(1), int64(13)
```

```
memory usage: 112.2 KB
```

```
df.describe()
```

	age	sex	cp	trestbps	chol
\count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000
std	9.072290	0.460373	1.029641	17.516718	51.59251
min	29.000000	0.000000	0.000000	94.000000	126.000000
25%	48.000000	0.000000	0.000000	120.000000	211.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000
75%	61.000000	1.000000	2.000000	140.000000	275.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000

	fbs	restecg	thalach	exang	oldpeak
\count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000

mean	0.149268	0.529756	149.114146	0.336585	1.071512
std	0.356527	0.527878	23.005724	0.472772	1.175053
min	0.000000	0.000000	71.000000	0.000000	0.000000
25%	0.000000	0.000000	132.000000	0.000000	0.000000
50%	0.000000	1.000000	152.000000	0.000000	0.800000
75%	0.000000	1.000000	166.000000	1.000000	1.800000
max	1.000000	2.000000	202.000000	1.000000	6.200000

	slope	ca	thal	target
count	1025.000000	1025.000000	1025.000000	1025.000000
mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

df.shape

(1025, 14)

Check for missing values in each column


df.isnull().sum()

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

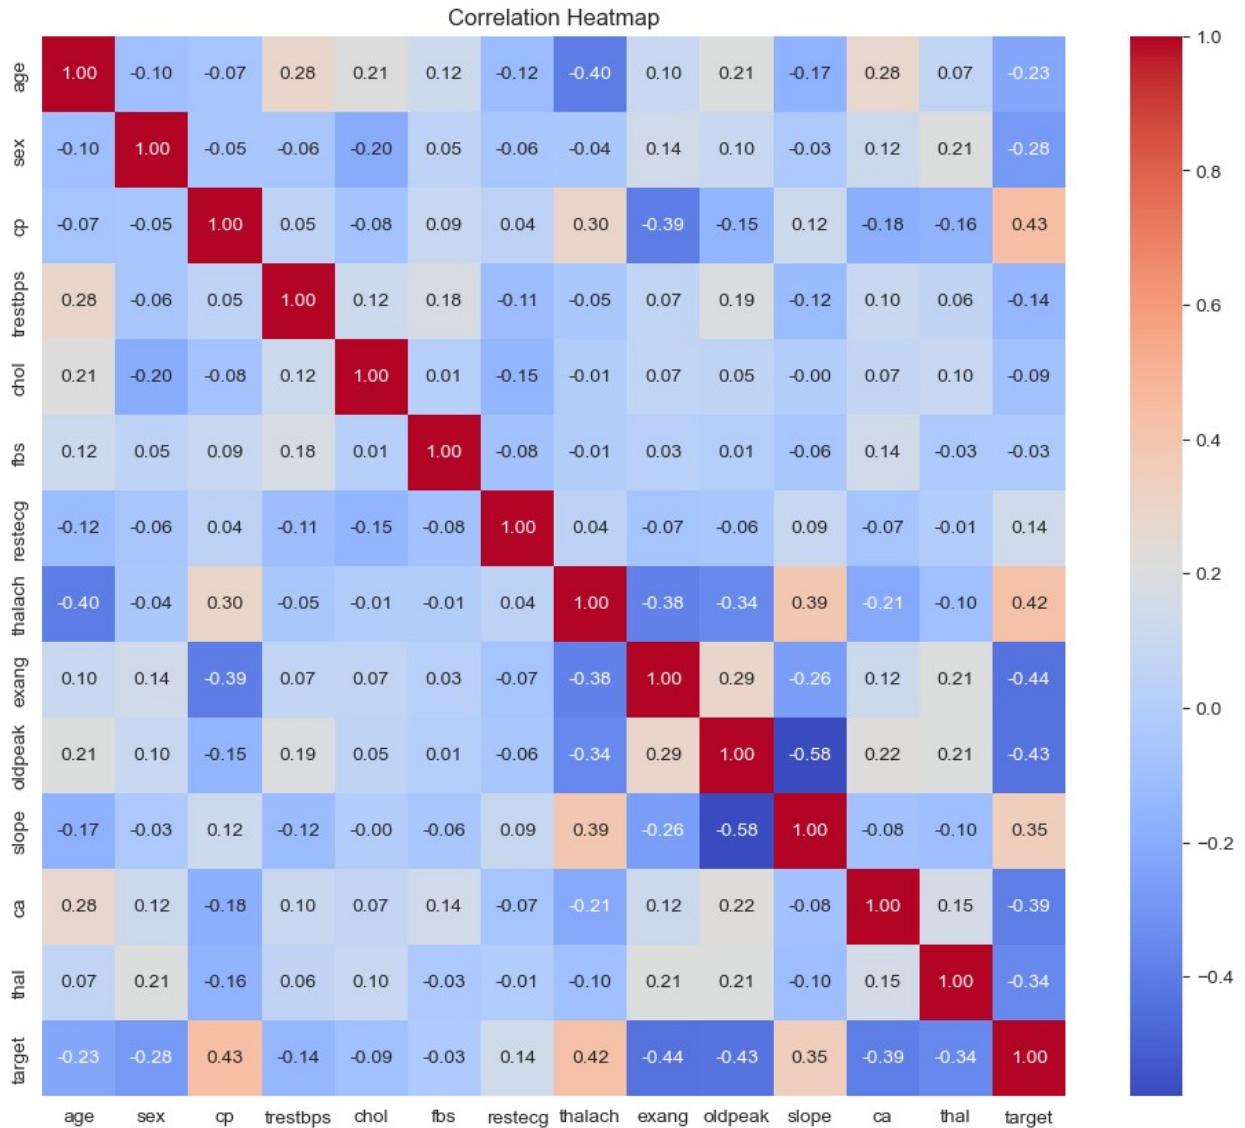
Plot histograms for all numerical variables to see their distributions

df.hist(figsize=(15, 15))

```
plt.tight_layout() # Adjust layout to prevent overlapping
plt.show()
```



```
numeric_cols = df.select_dtypes(include=['int64',
'float64']).columns.tolist()
# Plot heatmap for correlation matrix
plt.figure(figsize=(12,10))
corr = df[numeric_cols].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```



3.Data preprocessing

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
```

```
df.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach
exang \								
0	0.479167	1.0	0.0	0.292453	0.196347	0.0	0.5	0.740458
0.0								
1	0.500000	1.0	0.0	0.433962	0.175799	1.0	0.0	0.641221
1.0								
2	0.854167	1.0	0.0	0.481132	0.109589	0.0	0.5	0.412214

1.0								
3	0.666667	1.0	0.0	0.509434	0.175799	0.0	0.5	0.687023
0.0								
4	0.687500	0.0	0.0	0.415094	0.383562	1.0	0.5	0.267176
0.0								

	oldpeak	slope	ca	thal	target
0	0.161290	1.0	0.50	1.000000	0.0
1	0.500000	0.0	0.00	1.000000	0.0
2	0.419355	0.0	0.00	1.000000	0.0
3	0.000000	1.0	0.25	1.000000	0.0
4	0.306452	0.5	0.75	0.666667	0.0

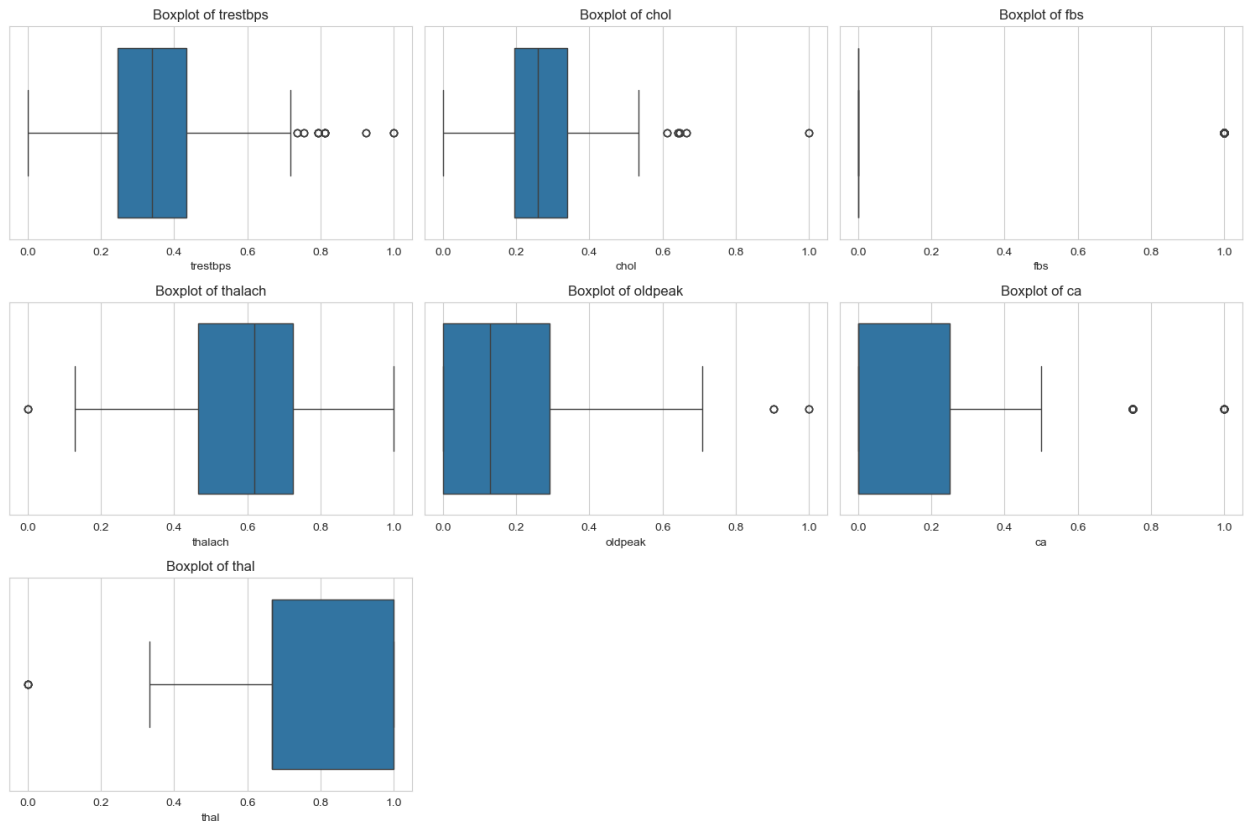
Milestone 2: Data Analysis and Visualization

1. Outliers detection and handling

```
import matplotlib.pyplot as plt
import seaborn as sns

# List of columns with outliers
cols_with_outliers = ['trestbps', 'chol', 'fbs', 'thalach', 'oldpeak',
                      'ca', 'thal']

# Plot boxplots for each column with outliers
plt.figure(figsize=(15, 10))
for i, col in enumerate(cols_with_outliers, 1):
    plt.subplot(3, 3, i)
    sns.boxplot(data=df, x=col)
    plt.title(f'Boxplot of {col}')
plt.tight_layout()
plt.show()
```



```
import pandas as pd

# List of columns with outliers to remove
cols_with_outliers = ['trestbps', 'chol', 'fbs', 'thalach', 'oldpeak',
                       'ca', 'thal']

def remove_outliers_iqr(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Keep rows within the bounds
    return df[(df[column] >= lower_bound) & (df[column] <=
upper_bound)]

df_clean = df.copy()

# Loop over columns and iteratively remove outliers
for col in cols_with_outliers:
    df_clean = remove_outliers_iqr(df_clean, col)

print(df_clean.skew(numeric_only=True))
```

```
age          -0.177371
sex          -0.832416
cp           0.580705
trestbps     0.276280
chol         0.266805
fbs          0.000000
restecg      -0.006654
thalach      -0.482755
exang        0.754599
oldpeak      0.944266
slope        -0.503237
ca           1.075802
thal         -0.143249
target       -0.196376
dtype: float64
```

```
import numpy as np
df_clean['ca' + '_log'] = np.log1p(df_clean['ca'])
print(df_clean.skew(numeric_only=True))
```

```
age          -0.177371
sex          -0.832416
cp           0.580705
trestbps     0.276280
chol         0.266805
fbs          0.000000
restecg      -0.006654
thalach      -0.482755
exang        0.754599
oldpeak      0.944266
slope        -0.503237
ca           1.075802
thal         -0.143249
target       -0.196376
ca_log       0.950673
dtype: float64
```

```
# Select numerical columns
numerical_cols = ['age', 'trestbps', 'chol', 'thalach']
```

```
# Plot histograms and KDE for each numerical column
```

```
plt.figure(figsize=(15, 10))
for i, col in enumerate(numerical_cols, 1):
    plt.subplot(2, 2, i)
```

```
# Histogram with KDE
```

```
sns.histplot(df_clean[col], kde=True, color='skyblue', bins=30)
plt.title(f'Distribution of {col}')
```



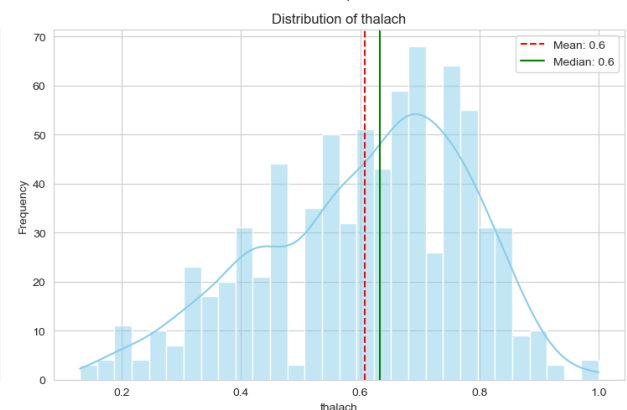
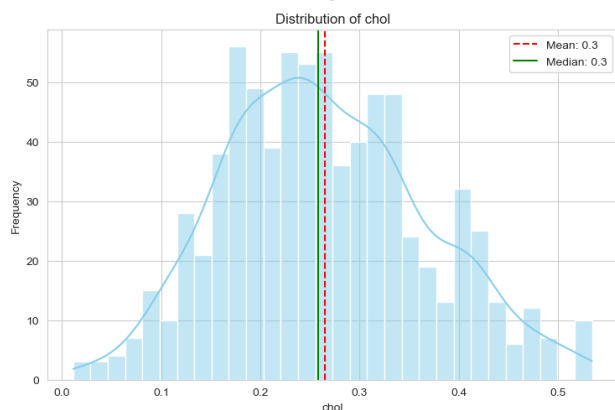
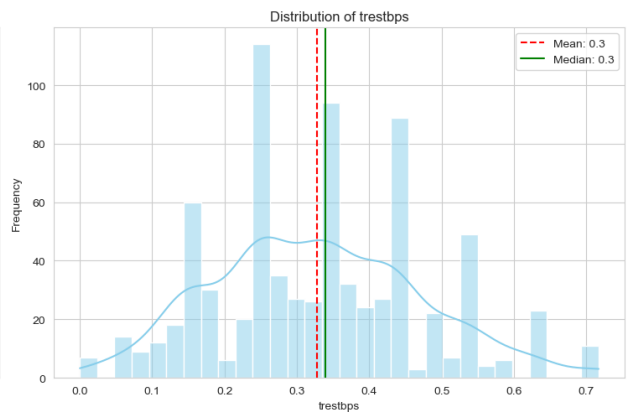

```

plt.xlabel(col)
plt.ylabel('Frequency')

# Mark mean/median for reference"
mean_val = df_clean[col].mean()
median_val = df_clean[col].median()
plt.axvline(mean_val, color='red', linestyle='--', label=f'Mean:
{mean_val:.1f}')
plt.axvline(median_val, color='green', linestyle='-',
label=f'Median: {median_val:.1f}')
plt.legend()

plt.tight_layout()
plt.show()

```



```
df_clean.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach
exang \								
0	0.479167	1.0	0.0	0.292453	0.196347	0.0	0.5	0.740458
0.0								
2	0.854167	1.0	0.0	0.481132	0.109589	0.0	0.5	0.412214
1.0								
3	0.666667	1.0	0.0	0.509434	0.175799	0.0	0.5	0.687023

```

0.0
5 0.604167 0.0 0.0 0.056604 0.278539 0.0 0.0 0.389313
0.0
7 0.541667 1.0 0.0 0.622642 0.372146 0.0 0.0 0.564885
1.0

```

	oldpeak	slope	ca	thal	target	ca_log
0	0.161290	1.0	0.50	1.000000	0.0	0.405465
2	0.419355	0.0	0.00	1.000000	0.0	0.000000
3	0.000000	1.0	0.25	1.000000	0.0	0.223144
5	0.161290	0.5	0.00	0.666667	1.0	0.000000
7	0.129032	0.5	0.25	1.000000	0.0	0.223144

Data Visualization


```

import matplotlib.pyplot as plt
import seaborn as sns

# Filter data for heart disease patients only
heart_disease_patients = df[df['target'] == 1]

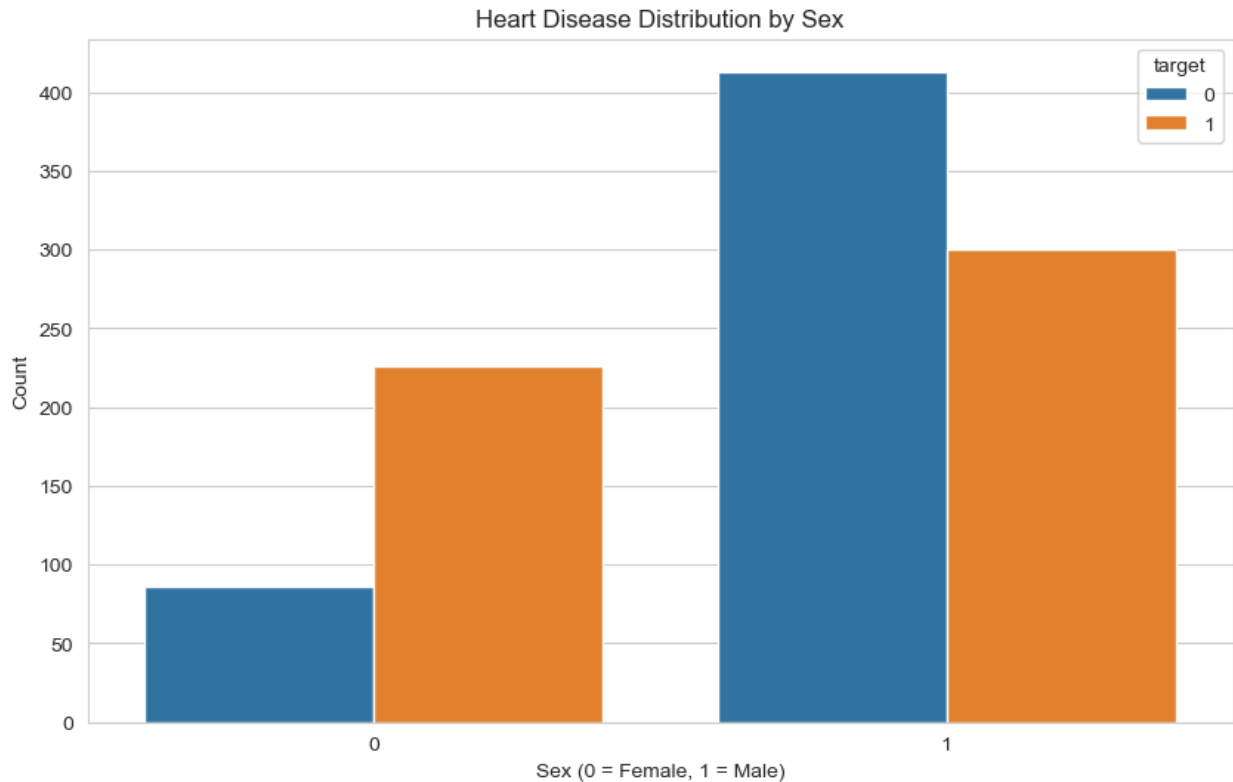
# Create histogram
plt.figure(figsize=(10, 6))
sns.histplot(data=heart_disease_patients, x='age', bins=15, kde=False,
color='red')
plt.title('Number of Heart Disease Patients by Age')
plt.xlabel('Age')
plt.ylabel('Number of Patients with Heart Disease')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()

```



```
import seaborn as sns
import matplotlib.pyplot as plt

# Plot count of target per sex
sns.countplot(x='sex', hue='target', data=df)
plt.title('Heart Disease Distribution by Sex')
plt.xlabel('Sex (0 = Female, 1 = Male)')
plt.ylabel('Count')
plt.show()
```



```
# Group data by sex and calculate the sum of target (heart disease
cases) for each sex
disease_counts = df.groupby('sex')['target'].sum()

# Print the number of females (sex=0) and males (sex=1) who have heart
disease
print("Number of females with heart disease:", disease_counts[0])
print("Number of males with heart disease:", disease_counts[1])

Number of females with heart disease: 226
Number of males with heart disease: 300

import matplotlib.pyplot as plt

# Count of target for each sex
female_counts = df[df['sex'] == 0]
['target'].value_counts().sort_index()
male_counts = df[df['sex'] == 1]['target'].value_counts().sort_index()

# Labels and colors
labels = ['No Disease', 'Has Disease']
colors = ['lightgreen', 'lightcoral']

# Create subplots
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
```


```

# Female pie chart
axes[0].pie(female_counts, labels=labels, autopct='%1.1f%%',
            colors=colors, startangle=90)
axes[0].set_title('Heart Disease in Females')

# Male pie chart
axes[1].pie(male_counts, labels=labels, autopct='%1.1f%%',
            colors=colors, startangle=90)
axes[1].set_title('Heart Disease in Males')

plt.tight_layout()
plt.show()

```



```

# Count samples per cp type
cp_counts = df['cp'].value_counts().sort_index()
print("Samples count per cp type:")
print(cp_counts)

# Calculate heart disease rate per cp type
cp_heart_disease_rate = df.groupby('cp')['target'].mean() * 100

print("\nHeart disease rate per cp type (%):")
print(cp_heart_disease_rate)

# Plot count of samples per cp type
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
sns.barplot(x=cp_counts.index, y=cp_counts.values, palette='Blues_d',
            hue=cp_counts.index, legend=False)

plt.title('Count of Samples by Chest Pain Type (cp)')
plt.xlabel('Chest Pain Type (cp)')

```

```
plt.ylabel('Number of Samples')

# Plot heart disease rate per cp type
plt.subplot(1, 2, 2)
sns.barplot(x=cp_counts.index, y=cp_counts.values, color='blue')
plt.title('Heart Disease Rate by Chest Pain Type')
plt.xlabel('Chest Pain Type (cp)')
plt.ylabel('Percentage with Heart Disease (%)')

plt.tight_layout()
plt.show()
```

Samples count per cp type:

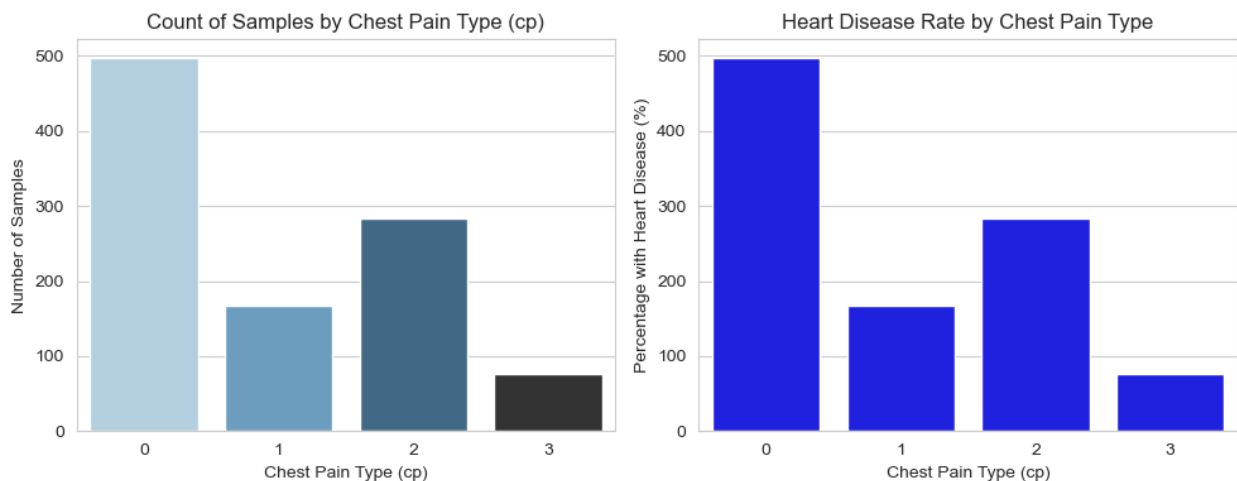
```
cp
0    497
1    167
2    284
3     77
```

Name: count, dtype: int64

Heart disease rate per cp type (%):

```
cp
0    24.547284
1    80.239521
2    77.112676
3    66.233766
```

Name: target, dtype: float64



```
import matplotlib.pyplot as plt
import seaborn as sns

# Calculate heart disease rate by exang
exang_counts = df['exang'].value_counts()
exang_heart_disease_rate = df.groupby('exang')['target'].mean() * 100
```

```

print("Samples count per exang value:")
print(exang_counts)

print("\nHeart disease rate by exang value (%):")
print(exang_heart_disease_rate)

# Plot for exang
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
sns.barplot(x=exang_counts.index, y=exang_counts.values,
palette='coolwarm')
plt.title('Count of Samples by exang (Exercise Induced Angina)')
plt.xlabel('exang (0 = No, 1 = Yes)')
plt.ylabel('Number of Samples')

plt.subplot(1, 2, 2)
sns.barplot(x=exang_heart_disease_rate.index,
y=exang_heart_disease_rate.values, palette='rocket')
plt.title('Heart Disease Rate by exang')
plt.xlabel('exang (0 = No, 1 = Yes)')
plt.ylabel('Percentage with Heart Disease (%)')

plt.tight_layout()
plt.show()

```

Samples count per exang value:

exang

0 680

1 345

Name: count, dtype: int64

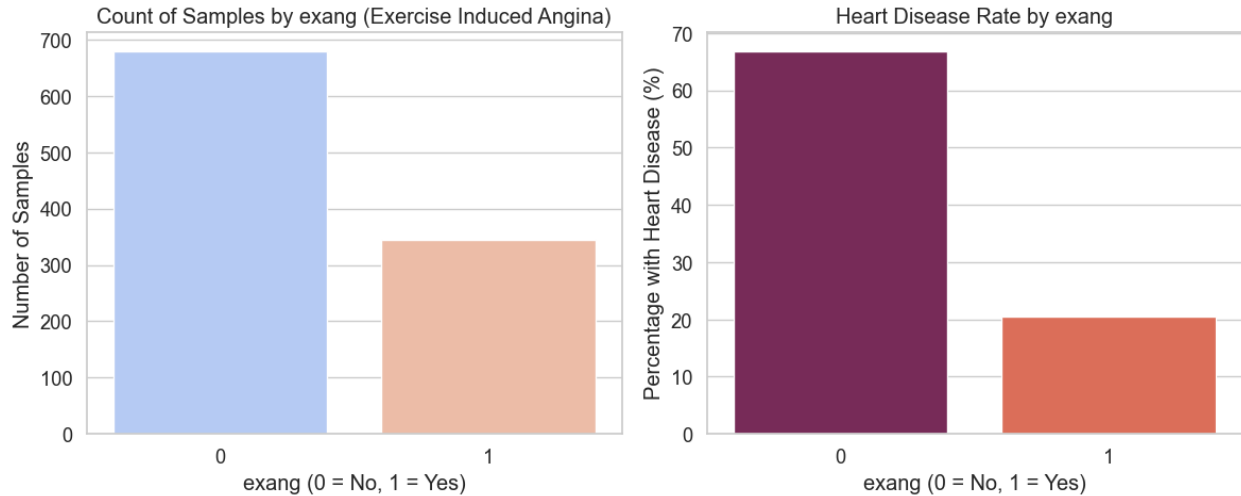
Heart disease rate by exang value (%):

exang

0 66.911765

1 20.579710

Name: target, dtype: float64







```
columns_to_plot = [ 'fbs', 'restecg', 'slope', 'ca', 'thal']


for col in columns_to_plot:
    plt.figure(figsize=(8,4))

    # Group by column value and sum target (number of patients with
    # heart disease)
    data_plot = df.groupby(col)['target'].sum().reset_index()

    sns.barplot(x=col, y='target', data=data_plot, hue=col,
    palette='viridis', legend=False)
    plt.title(f'Number of Heart Disease Patients by {col}')
    plt.ylabel('Number of Patients with Heart Disease')
    plt.xlabel(col)
    plt.show()
```





Age group dashboard

```
import pandas as pd
from dash import Dash, dcc, html
import plotly.express as px

# Load your dataset CSV file (adjust path as needed)
df = pd.read_csv('heart.csv')

# Example: If you don't have an 'age_group' column,
# create it by binning the 'age' column
bins = [25, 34, 44, 54, 64, 79]
labels = ['25-34', '35-44', '45-54', '55-64', '65-79']
df['age_group'] = pd.cut(df['age'], bins=bins, labels=labels,
right=True, include_lowest=True)

# Count number of patients per age group
age_group_counts = df['age_group'].value_counts().sort_index()
age_group_df = age_group_counts.reset_index()
age_group_df.columns = ['age_group', 'patient_count']

# Create the Dash app
app = Dash(__name__)

# Create bar chart of patient count by age group
fig = px.bar(
    age_group_df,
    x='age_group',
    y='patient_count',
```

```

        title='Number of Patients per Age Group',
        labels={'age_group': 'Age Group', 'patient_count': 'Number of
Patients'},
        template='plotly_white'
    )

# Layout
app.layout = html.Div(children=[
    html.H1('Patient Count by Age Group Dashboard'),
    dcc.Graph(id='patient-count-bar', figure=fig)
])

if __name__ == '__main__':
    app.run_server(debug=True)

<IPython.lib.display.IFrame at 0x2520803de20>

```

Milestone 3: Predictive Model Development and Optimization

```

# Calculate correlation of all columns with 'target'
correlations = df_clean.corr()['target'].abs()

# Filter correlations with absolute value greater than 0.3
filtered_corr = correlations[correlations.abs() > 0.3]

# Sort the filtered correlations by absolute value in descending order
filtered_corr = filtered_corr.sort_values(key=abs, ascending=False)

# Print the filtered correlations
print(filtered_corr)

```

target	1.000000
ca_log	0.464166
ca	0.456208
thal	0.454862
oldpeak	0.443599
thalach	0.416945
exang	0.416859
cp	0.374184
sex	0.367541
slope	0.319420

Name: target, dtype: float64

1. Model Training and evaluation:

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score

```

```

from sklearn.preprocessing import StandardScaler

# Step 1: Select features and target
features = ['ca_log', 'thal', 'oldpeak', 'thalach', 'exang', 'cp',
'sex', 'slope']
X = df_clean[features]
y = df_clean['target']

# Step 2: Normalize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Split dataset into train and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42)

# Step 4: Build Logistic Regression model
model = LogisticRegression(random_state=42)
model.fit(X_train, y_train)

# Step 5: Predict on test set
y_pred = model.predict(X_test)

# Step 6: Evaluate the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test,
y_pred))

y_train_pred = model.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)

print("Train Accuracy:", train_accuracy)

Accuracy: 0.8571428571428571

Classification Report:

```

	precision	recall	f1-score	support
0.0	0.86	0.78	0.82	64
1.0	0.85	0.91	0.88	90
accuracy			0.86	154
macro avg	0.86	0.85	0.85	154
weighted avg	0.86	0.86	0.86	154

```

Train Accuracy: 0.8666666666666667

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix


```

```

# Compute the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Plot the confusion matrix as a heatmap
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=[0, 1],
            yticklabels=[0, 1])
plt.xlabel('Predicted Label') # Label for x-axis
plt.ylabel('True Label')      # Label for y-axis
plt.title('Confusion Matrix') # Title of the plot
plt.show()

```



Milestone 4: Deployment

```

import joblib

# Save the trained model
joblib.dump(model, 'heart_disease_model.pkl')

['heart_disease_model.pkl']

import streamlit

!streamlit run app.py

```